DMMST — Version 2 results: our loss variants compared with other models

Parsa

1 Setup

We compare our model (a transformer with the survival head of the paper) trained with different losses against eight published survival models. All models are run on the same five datasets under the same protocol:

- 10 random splits of each dataset (60% train, 10% validation, 30% test); results are reported as mean (standard deviation) over the 10 splits;
- settings of every model chosen on the validation split only;
- every model scored by the same code.

Datasets. METABRIC and SUPPORT (real, one event: death); EBMT (real, 4 events); hsa synthetic and DeepHit synthetic (published synthetic data, 2 events). SEER was planned as an additional large real dataset, but because of the updated SEER+ data-access policies we could not obtain access to it.

Our variants.

Name	Losses
S1	L_{PCH} (survival likelihood)
S2	$L_{PCH} + L_{rank}$
S3	$L_{PCH} + L_{mul}$ (multi-event data only)
S4	$L_{PCH} + L_{rank} + L_{mul}$ (multi-event data only)
Full	survival head with all its losses + regression head (L_{MAE} with best guess)

L_{mul} compares different events of the same patient, so S3 and S4 do not exist on single-event data ("—" in the tables).

Regression head. Besides the survival curve, the model can predict one number: the time of the event. For censored patients the true time is unknown; R1 simply leaves them out, the others use a "best guess" (from a Kaplan–Meier curve of the training data) as the target.

Name	Losses
R1	L_{MAE} on observed events only (censored patients ignored)
R2	L_{MAE} with the best guess for censored patients
R3	R2 + L_{MM} (mismatch loss), observed times (multi-event data only)
R4	R2 + L_{MM}, best guesses (multi-event data only)
Joint	both heads trained together, every S × R combination

Predicted times are scored by MAE (mean absolute error, lower is better; in the time unit of each dataset) and Harrell C (how well predicted times rank the patients, higher is better).

Baselines. SurvTRACE (authors' code), DeepHit, RSF, PC-Hazard, DeepSurv, MENSA, Cox, DSM.

Scores. C_{td}: how well patients are ranked by risk (higher is better; 0.5 is random). IBS: integrated Brier score, the accuracy of the predicted survival probabilities (lower is better).

2 Results

Table 1: C_{td}, mean (sd) over 10 splits. Best per dataset in bold.

Model	METABRIC	SUPPORT	EBMT	hsa synth.	DeepHit synth.
Ours S1	0.667 (0.013)	0.615 (0.011)	0.529 (0.013)	0.849 (0.049)	0.743 (0.003)
Ours S2	0.670 (0.016)	0.620 (0.007)	0.535 (0.016)	**0.868** (0.010)	0.743 (0.004)
Ours S3	—	—	0.530 (0.016)	0.859 (0.015)	0.742 (0.003)
Ours S4	—	—	0.537 (0.015)	0.867 (0.009)	0.742 (0.004)
Ours Full	0.672 (0.021)	0.619 (0.007)	0.534 (0.017)	0.852 (0.041)	0.744 (0.004)
SurvTRACE	**0.678** (0.013)	**0.638** (0.008)	0.546 (0.009)	**0.868** (0.008)	0.755 (0.007)
DeepHit	0.677 (0.012)	0.610 (0.008)	0.525 (0.011)	0.865 (0.010)	0.737 (0.006)
RSF	0.672 (0.014)	0.636 (0.007)	**0.549** (0.009)	0.849 (0.009)	0.733 (0.005)
PC-Hazard	0.668 (0.017)	0.627 (0.008)	0.548 (0.010)	0.837 (0.008)	**0.762** (0.003)
DeepSurv	0.640 (0.030)	0.593 (0.007)	0.529 (0.014)	0.842 (0.009)	**0.762** (0.002)
MENSA	0.654 (0.017)	0.596 (0.021)	0.501 (0.019)	0.847 (0.014)	0.570 (0.074)
Cox	0.631 (0.022)	0.567 (0.006)	0.531 (0.014)	0.497 (0.013)	0.591 (0.005)
DSM	0.620 (0.065)	0.602 (0.037)	0.512 (0.013)	0.522 (0.031)	0.653 (0.104)

Table 2: Integrated Brier score (lower is better), mean (sd) over 10 splits. Best per dataset in bold.

Model	METABRIC	SUPPORT	EBMT	hsa synth.	DeepHit synth.
Ours S1	0.179 (0.012)	0.200 (0.004)	0.338 (0.165)	0.075 (0.007)	0.174 (0.005)
Ours S2	0.175 (0.006)	0.199 (0.004)	0.313 (0.153)	0.075 (0.003)	0.171 (0.005)
Ours S3	—	—	0.256 (0.082)	0.074 (0.003)	0.175 (0.005)
Ours S4	—	—	0.233 (0.043)	0.075 (0.003)	0.172 (0.006)
Ours Full	0.177 (0.005)	0.200 (0.003)	0.299 (0.138)	0.076 (0.005)	0.172 (0.007)
SurvTRACE	0.201 (0.033)	**0.193** (0.004)	0.242 (0.018)	**0.071** (0.003)	0.169 (0.006)
DeepHit	0.198 (0.007)	0.217 (0.004)	0.388 (0.159)	0.091 (0.002)	0.254 (0.012)
RSF	0.176 (0.007)	**0.193** (0.003)	0.224 (0.009)	0.087 (0.002)	0.185 (0.005)
PC-Hazard	**0.174** (0.006)	0.194 (0.004)	0.219 (0.010)	0.078 (0.003)	0.171 (0.005)
DeepSurv	0.176 (0.006)	0.195 (0.004)	0.220 (0.008)	0.078 (0.003)	**0.167** (0.005)
MENSA	0.189 (0.012)	0.212 (0.014)	0.331 (0.023)	0.079 (0.003)	0.223 (0.014)
Cox	0.175 (0.006)	0.198 (0.004)	**0.218** (0.009)	0.105 (0.002)	0.218 (0.006)
DSM	0.224 (0.067)	0.205 (0.008)	0.326 (0.024)	0.369 (0.049)	0.232 (0.078)

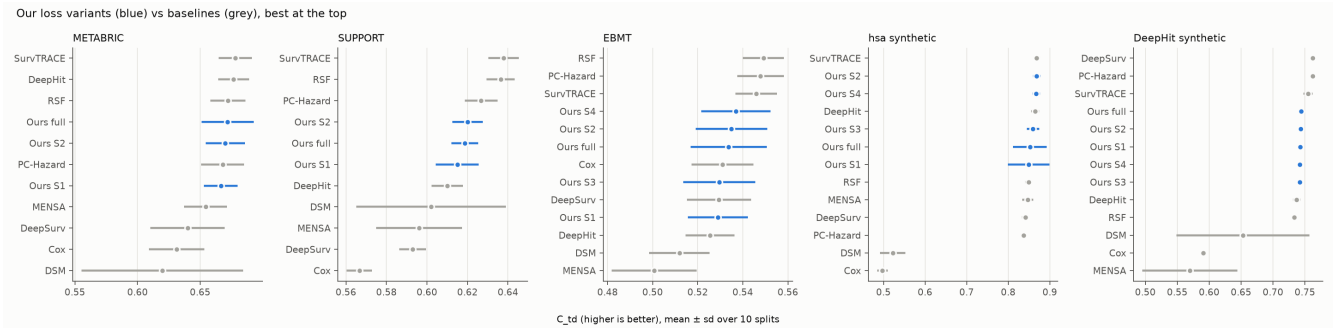


Figure 1: C_{td} of our variants (blue) and the baselines (grey), best at the top of each panel.

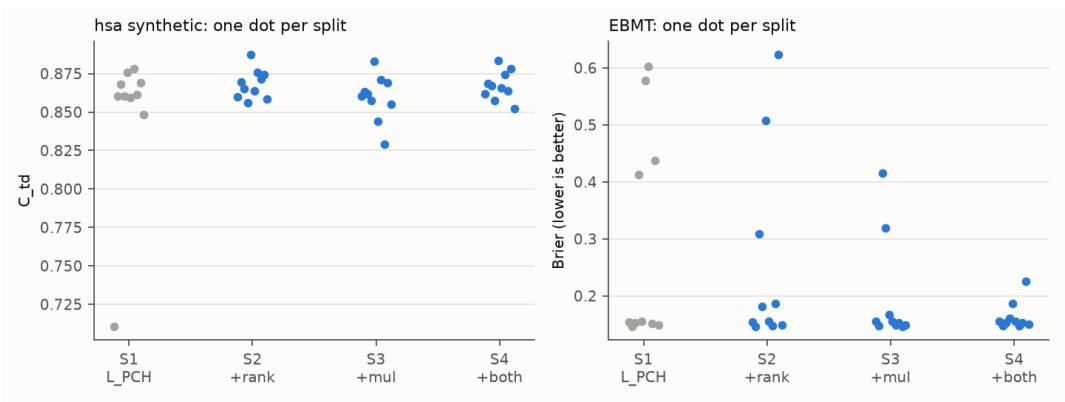


Figure 2: Every split of each loss variant. Left: C_{td} on hsa synthetic. Right: Brier score on EBMT. With L_{PCH} alone (S1) some splits fail; with $L_{rank} + L_{mul}$ (S4) none do.

Regression head

Table 3: Predicted times: Harrell C (higher is better), mean (sd) over 10 splits. Best per dataset in bold. "Joint" is the mean over all S × R combinations (both heads trained together). n/a: L_{MM} needs several events, so R3/R4 do not exist on METABRIC and SUPPORT.

Recipe	METABRIC	SUPPORT	EBMT	hsa synth.	DeepHit synth.
R1	0.608 (0.018)	0.546 (0.007)	0.510 (0.012)	0.566 (0.021)	0.689 (0.046)
R2	**0.643** (0.012)	**0.592** (0.009)	0.524 (0.014)	0.663 (0.047)	0.704 (0.042)
R3	n/a	n/a	0.530 (0.009)	0.663 (0.087)	0.700 (0.035)
R4	n/a	n/a	**0.533** (0.011)	0.665 (0.063)	**0.714** (0.003)
Joint	0.626 (0.026)	0.575 (0.023)	0.504 (0.015)	**0.689** (0.125)	0.551 (0.141)

Table 4: Predicted times: MAE (lower is better), mean (sd) over 10 splits, in the time unit of each dataset (so values can only be compared within a column). n/a as in Table 3.

Recipe	METABRIC	SUPPORT	EBMT	hsa synth.	DeepHit synth.
R1	80.6 (2.8)	613.9 (9.1)	4254 (123)	396.4 (5.8)	38.8 (1.5)
R2	**76.7** (3.2)	**604.6** (9.6)	**4232** (130)	408.7 (96.7)	34.8 (3.2)
R3	n/a	n/a	4245 (126)	420.6 (94.9)	35.0 (3.0)
R4	n/a	n/a	4247 (120)	**379.8** (62.2)	**33.9** (0.4)
Joint	80.7 (4.7)	612.3 (13.7)	4321 (150)	487.5 (120.6)	44.3 (5.6)

What these tables show.

- **Using the best guess for censored patients (R2) is better than ignoring them (R1).** R2 ranks predicted times better on every dataset (Harrell C 0.566 $\rightarrow$ 0.663 on hsa synthetic, 0.546 $\rightarrow$ 0.592 on SUPPORT) and has a lower MAE on four of five datasets (hsa synthetic is the exception).
- **The mismatch loss (R3, R4) adds little.** R4 is the best recipe on EBMT, hsa and DeepHit synthetic, but only by small margins over R2.
- **Joint training:** on hsa synthetic training both heads together ranks times better than the regression head alone (0.689 vs about 0.66). On DeepHit synthetic it is much worse (0.551 vs about 0.70), because joint models pick their checkpoint by the survival C_{td}, not by the quality of the predicted times. The large sd (0.125, 0.141) shows that the result depends strongly on the recipe.
- **EBMT:** Harrell $C \approx 0.5$ in every setup, so predicted times there carry almost no information (the same as for the survival head).

3 Findings

1. **Our model is competitive but not the best.** It is in the upper half on every dataset and ties SurvTRACE on hsa synthetic, but wins no dataset. Our best variant is 0.000–0.018 C_{td} behind the best model.
2. L_{rank} **helps a little on every dataset** (S2 $\geq$ S1). The gain is statistically significant only on hsa synthetic.
3. **The extra losses make training reliable.** With L_{PCH} alone one hsa split collapses and several EBMT splits are badly calibrated. With $L_{rank} + L_{mul}$ (S4) this does not happen: the EBMT IBS drops from 0.338 to 0.233.
4. **Adding the regression head does not hurt survival prediction** when L_{rank} is used ("Full" $\approx$ S2/S4).
5. **For the regression head the best guess matters most** (R2 better than R1); L_{MM} adds little.
6. **Joint training helps predicted times on hsa synthetic but hurts them on DeepHit synthetic** (Harrell C 0.689 vs 0.551), because the checkpoint is chosen by the survival C_{td}, not by the regression score.
7. **Calibration is good** on four of five datasets, and better than SurvTRACE on METABRIC. EBMT is the exception.
8. **EBMT is hard for every model** (C_{td} 0.50–0.55): it has only four categorical features.

4 Conclusion

The ranking losses of the paper do not make the model more accurate than the strongest baselines, but they make it more reliable and better calibrated, and they allow the regression head to be trained together with the survival head at no cost. The results do not support a claim of better C_{td} than SurvTRACE.